

# Python for HPC

**Andrea Zonca - SDSC**

# Jupyter Notebook

**Data exploration in your browser**

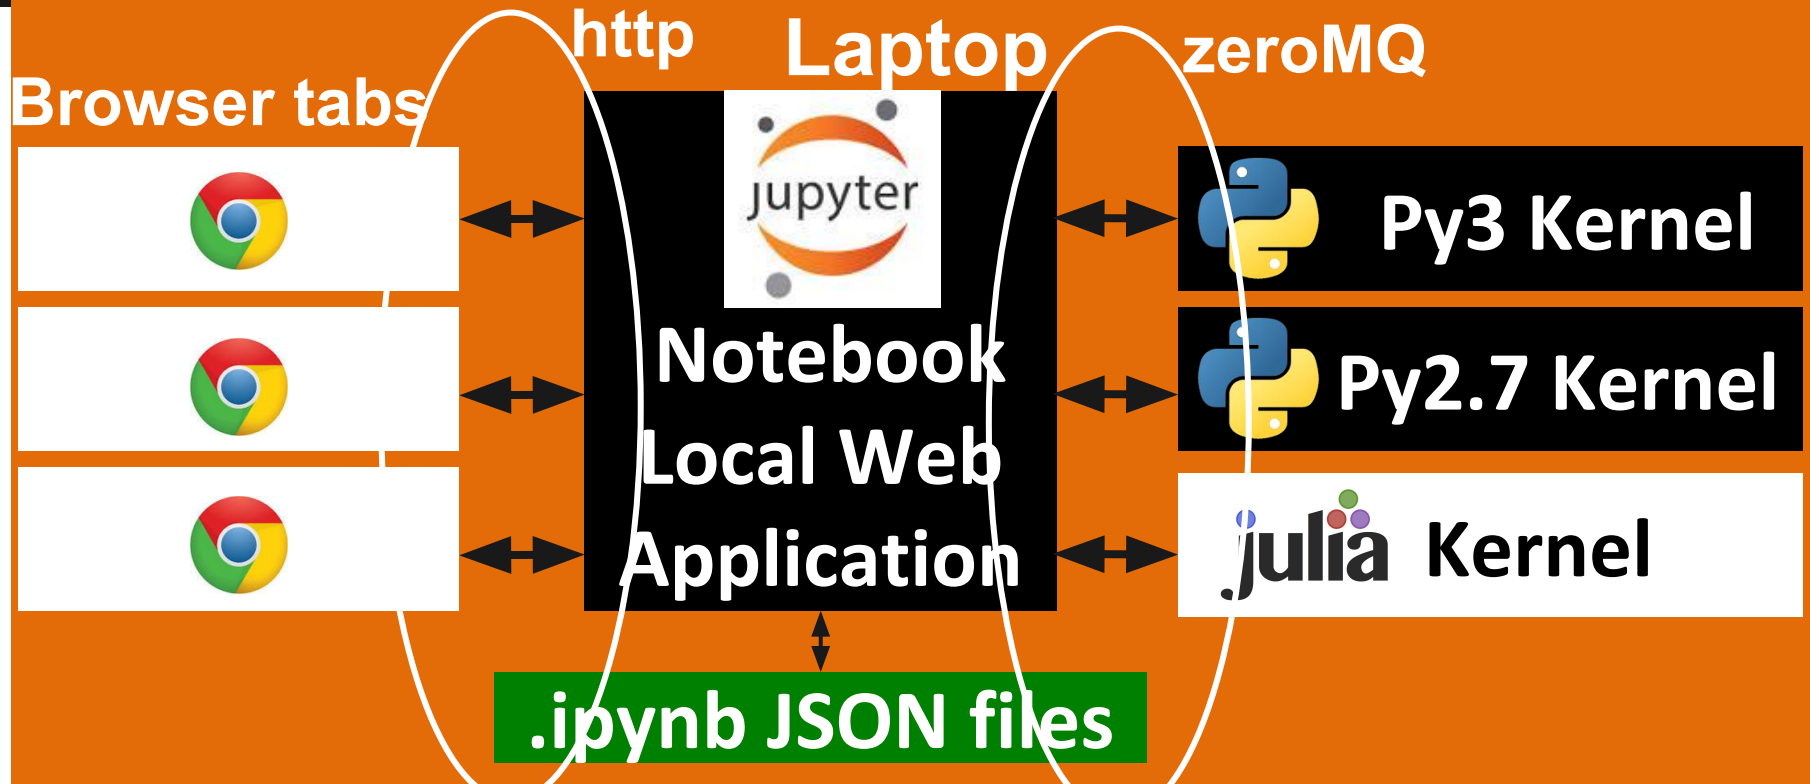
# What is the notebook?

- Browser based interactive console
- Supports multiple sessions in browser tabs
- Each session has a Kernel executing computation
- Saved in JSON format

# Notebooks on Nature

<http://www.nature.com/news/interactive-notebooks-sharing-the-code-1.16261>

# Jupyter notebook local



# Jupyter notebook remote

Laptop



https +  
password

Server

Jupyter  
Notebook  
Web  
Application



Py3 Kernel



Py2.7 Kernel



Kernel

.ipynb JSON files

# Clone workshop repository

ssh into comet with training account

```
git clone URL workshop
```

URL is <https://github.com/sdsc-scicomp/2015-11-12-ucla>

# Modules on Comet

```
module load python scipy
```

- add line to `.bashrc`



# Setup on Comet

- ssh to Comet

```
salloc --nodes=1 --tasks-per-node=1 -t 02:00:00 --res=ucla2015
```

- `ssh comet-xx-xx`
- `ipython notebook --no-browser --ip="*" &`
- later better setup config file
-

# connect with browser

Open browser on your laptop and connect to  
comet-xx-xx.sdsc.edu:8888

New -> Notebook

!hostname

# More secure setup

<http://zonca.github.io/2015/09/ipython-jupyter-notebook-sdsc-comet.html>

# IPython notebook demo

- Python code
- Formatted text
- Equations
- Plots
- Cells execution, cells order
- Clear output

# Why the notebook?

- Literate programming: code and explanation together
- Reproducible science: document easily every step
- Easy to share computations: send one single notebook instead of scripts/plots/.doc

# ipynb documents

- JSON format
- includes plots in binary format
- easy to convert to .html/.pdf for sharing
- <http://nbviewer.ipython.org>
- Recently rendered automatically on Github

# HPC: interactive notebooks

- Analyze large amount of data
- In-situ visualization
- Centralized Python stack
- Check long-running computations
- Prepare and submit batch jobs

# Notebooks as scripts

- Install runipy:

```
pip install --user runipy
```

- Setup .bashrc:

```
cd ~/workshop/python_hpc
```

```
cat setup_pip_local.sh >> ~/.bashrc
```

- Restart bash with: `bash`



# Notebooks as scripts

- demo of runipy
  - open and execute `fit_line.ipynb`
  - uncomment cell with `(os.environ)`
  - `white_noise_scale=1000 runipy fit_line.ipynb fit_line_1000.ipynb`
  - open `fit_line_1000.ipynb`, what happened?
- demo of batch submission of SLURM serial runipy jobs using pipes

# Hands-on

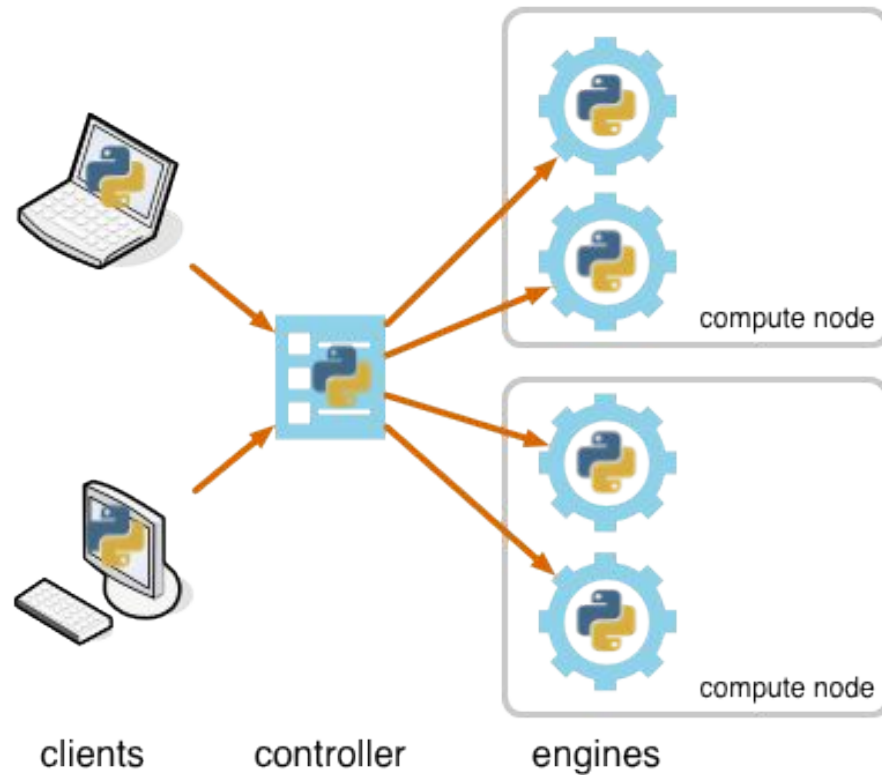
- Open the notebook interactively
- Add saving the plot with `plt.savefig("figurename.png")` in the same cell
- Test with `runipy` on the interactive node
- Rerun the jobs through the queue

# IPython parallel

Parallel computing the easy way

# IPython parallel

- High-level API for distributed computing with Python
- Engines (Python worker processes) connected to Controller with ZeroMQ
- Client, i.e. user's IPython session, connects to the Controller



*independent python kernel*

Image by Continuum Analytics

# IPython parallel architecture

# Functionalities

- Load balanced queue for trivially parallel jobs
- Supports job dependencies
- Direct interface to Engines
- Supports MPI applications, Python or C/C++/Fortran

# IPython parallel config

```
ipython profile create
```

```
cp ipython_parallel_configuration/* ~/.  
ipython/profile_default
```

# IPython parallel Demo

- Launch cluster with 48 engines:
  - `ipcluster start --n=48`
- Connect with IPython Notebook
- Print ids, hostnames
- Launch demo job and check it runs correctly



# Hands-on

- Create a duplicate of `fit_line.ipynb`
- Reformat `fit_line` code into a single function
- Send it to engines for execution within the balanced queue
- Print out the results from the notebook



```
In [1]: from IPython.parallel import Client
```

```
In [2]: c = Client()
```

```
In [3]: view = c[:]
```

```
In [4]: view.activate() # enable magics
```

```
# run the contents of the file on each engine:
```

```
In [5]: view.run('psum.py')
```

```
In [6]: view.scatter('a', np.arange(16, dtype='float'))
```

```
In [7]: view['a']
```

```
Out[7]: [array([ 0.,  1.,  2.,  3.]),  
         array([ 4.,  5.,  6.,  7.]),  
         array([ 8.,  9., 10., 11.]),  
         array([12., 13., 14., 15.])]
```

```
In [7]: %px totalsum = psum(a)
```

```
Parallel execution on engines: [0,1,2,3]
```

```
In [8]: view['totalsum']
```

```
Out[8]: [120.0, 120.0, 120.0, 120.0]
```